

4703 Case Study

April 24, 2022

1 IEOR4703 Monte Carlo Case Study - 2 May 2022

Chun Yat Yeung (cy2623)

2 Setup

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
from scipy.stats import norm, t, gaussian_kde
```

```
[2]: df = pd.read_csv('20_funds.csv', index_col=0)
df.index = pd.date_range(start='12-31-2005', periods=len(df))
funds = df.columns
df.head()
```

```
[2]:
```

	B23409	B00367	B13847	B15553	B10749	B03605	\
2005-12-31	12.88289	35.55018	11.50636	13.29872	25.17435	19.75900	
2006-01-01	12.88289	35.55018	11.50636	13.29872	25.17435	19.75900	
2006-01-02	12.88289	35.55018	11.50636	13.29872	25.17435	19.75900	
2006-01-03	13.08824	35.60458	11.57815	13.52239	25.56893	20.38565	
2006-01-04	13.17470	35.63178	11.58840	13.60373	25.63469	20.69897	


	B17517	B05851	B12869	B20167	B18365	B21754	\
2005-12-31	10.44675	30.24456	10.00	246.50262	51.20245	28.29245	
2006-01-01	10.44675	30.24456	10.00	246.50262	51.20245	28.29245	
2006-01-02	10.44675	30.24456	10.00	246.50262	51.20245	28.29245	
2006-01-03	10.75808	30.26963	10.16	235.42289	51.59738	28.32407	
2006-01-04	10.89644	30.30422	10.23	234.30632	51.67636	28.32749	

	B21401	B21937	B21046	B18646	B11920	B20159	\
2005-12-31	71.00312	13.33306	21.25487	12.79461	18.93086	156.02005	
2006-01-01	71.00312	13.33306	21.25487	12.79461	18.93086	156.02005	
2006-01-02	71.00312	13.33306	21.25487	12.79461	18.93086	156.02005	
2006-01-03	71.96846	13.52990	21.32307	12.95087	18.98386	156.34645	
2006-01-04	72.50194	13.58170	21.37990	13.02773	19.07706	155.77525	

	B07325	B21967
2005-12-31	23.07846	25.53953
2006-01-01	23.07846	25.53953
2006-01-02	23.07846	25.53953
2006-01-03	23.76154	26.36954
2006-01-04	24.02989	26.80495

3 Visualization

```
[3]: n = 5
for fund in funds[:n]:
    plt.plot(df[fund],label=fund)
plt.xlabel('time')
plt.ylabel('fund NAV')
plt.legend()
plt.show()
```



4 Functions

```
[4]: def calcHistoricalReturns(df, rollingPeriod=30, timeHorizon=30, startDate=None,
    ↪endDate=None):
    ''' calculate historical returns of all mutual funds '''
    if not startDate: startDate = df.index[0]
    if not endDate: endDate = df.index[-1]
    rollIdx = np.arange(0, len(df), rollingPeriod)
    dfrtn = (df.shift(-timeHorizon) - df) / df
    dfrtn = dfrtn.iloc[rollIdx].loc[startDate:endDate]
    return dfrtn

def calcVaR(dfrtn, funds=None, allocation=None, confidence=0.99, Nsim=1e5,
    ↪method='bootstrap'):
    ''' calculate VaR & CVaR based on dfrtn with user-defined allocation/
    ↪confidence/Nsim/method '''
    if not funds: funds = dfrtn.columns # default: all funds
    if not allocation: allocation = np.repeat(1/len(funds), len(funds)) #
    ↪default: uniform
    portRtn = dfrtn[funds].dot(allocation).dropna().to_numpy() # historical
    ↪portfolio returns
    out = {'funds': funds, 'allocation': allocation, 'portRtn': portRtn}
    VaR = CVaR = None
    if method == 'bootstrap':
        # the prescribed approach in instructions
        # 1. construct historical portfolio returns, via: portRtn
        # 2. bootstrap-simulate VaR & CVaR, via: varSim, cvarSim
        Nsim = int(Nsim)
        varSim = np.zeros(Nsim)
        cvarSim = np.zeros(Nsim)
        varIdx = np.ceil(len(portRtn)*(1-confidence)).astype('int')
        for i in tqdm(range(Nsim)):
            bootRtn = np.sort(np.random.choice(portRtn, len(portRtn)))
            varSim[i] = bootRtn[varIdx]
            cvarSim[i] = np.mean(bootRtn[:varIdx])
        VaR = np.mean(varSim)
        CVaR = np.mean(cvarSim)
    elif method == 'historical':
        portRtnSort = np.sort(portRtn)
        varIdx = np.floor(len(portRtn)*(1-confidence)).astype('int')
        VaR = portRtnSort[varIdx]
        CVaR = np.mean(portRtnSort[:varIdx])
    elif method == 'normal':
        portRtnMu, portRtnSig = norm.fit(portRtn)
        VaR = norm.ppf(1-confidence, loc=portRtnMu, scale=portRtnSig)
        CVaR = norm.expect(lambda x: x, loc=portRtnMu, scale=portRtnSig,
    ↪ub=VaR, conditional=True)
```

```

elif method == 't-dist':
    portRtnDf, portRtnMu, portRtnSig = t.fit(portRtn)
    VaR = t.ppf(1-confidence, portRtnDf, loc=portRtnMu, scale=portRtnSig)
    CVaR = t.expect(lambda x: x, (portRtnDf,), loc=portRtnMu,
→scale=portRtnSig, ub=VaR, conditional=True)
    VaR, CVaR = -VaR, -CVaR
    out.update({'VaR': VaR, 'CVaR': CVaR})
    return out

def histBootVaR(df, timeHorizon=30, lookback=None, funds=None, allocation=None,
→confidence=0.99, Nsim=1e5):
    ''' alternative approach to calculating VaR & CVaR, by simulating entire
→paths of mutual funds, known as historical bootstrap '''
    if not funds: funds = df.columns # default: all funds
    if not allocation: allocation = np.repeat(1/len(funds), len(funds)) #
→default: uniform
    dfrtn = (df.shift(-1)-df)/df
    dfrtn = dfrtn[funds]
    if lookback: dfrtn = dfrtn.iloc[-lookback:]
    Nsim = int(Nsim)
    portNav = np.zeros(Nsim)
    for i in tqdm(range(Nsim)):
        portNav[i] = (dfrtn.iloc[np.random.choice(len(dfrtn), timeHorizon)]+1).
→product().dot(allocation)
    portRtn = portNav-1
    portRtnSort = np.sort(portRtn)
    varIdx = np.floor(len(portRtn)*(1-confidence)).astype('int')
    VaR = portRtnSort[varIdx]
    CVaR = np.mean(portRtnSort[:varIdx])
    VaR, CVaR = -VaR, -CVaR
    out = {'VaR': VaR, 'CVaR': CVaR}
    return out

```

5 Analysis

5.1 Bootstrap portfolio returns

We follow the below steps, as prescribed in instructions:

1. calculate historical **mutual fund** returns, based on user-defined `rollingPeriod` and `timeHorizon`
2. calculate historical **portfolio** returns, constructed with specified `funds` and `allocation`
3. **bootstrap-simulate** from historical portfolio returns (the empirical distribution), which leads to `varSim` and `cvarSim`, and the means constitute **VaR** and **CVaR**

`calcVaR(dfrtn, ...)` handles this routine, and offers a few other VaR calculation methods for comparison, including `historical`, `normal` and `t-dist`.

1. **historical** directly gives VaR and CVaR from historical portfolio returns, by sorting and finding the appropriate threshold
2. **normal** assumes a Gaussian distribution for portfolio returns, matching first and second moments, and this is known to severely underestimate tail risks, with very small VaR and CVaR
3. **t-dist** assumes a Student's t-distribution for portfolio returns, with a degree of freedom calibrated to historical returns. While it produces heavy tails, it assumes symmetric distribution, which may not be realistic (loss distribution is right-skewed)

In the example below, we study calculations for `rollingPeriod=30` and `timeHorizon=30`, i.e. non-overlapping **one-month** rolling returns. First obtain `dfrtn` from `calcHistoricalReturns(df, ...)`, then pass `dfrtn` to `calcVaR(dfrtn, ...)` with a desired method. With bootstrap simulation, we obtain the one-month VaR 0.12233% and CVaR 0.15118%. Note that these refer to portfolio **loss** over a one-month period. Roughly they have identical order of magnitude with **historical** method, but **historical** makes use of only one sample series, thus the estimated VaR and CVaR are not statistically significant. In contrast, **bootstrap** uses `Nsim` series, with estimations more stable (due to averaging over all bootstrap samples).

```
[5]: dfrtn = calcHistoricalReturns(df, 30, 30)
     dfrtn.head()
```

```
[5]:
```

	B23409	B00367	B13847	B15553	B10749	B03605	\
2005-12-31	0.046141	0.003443	0.015152	0.035933	0.038140	0.065411	
2006-01-30	0.005613	0.002669	0.014048	0.006642	-0.001007	0.007442	
2006-03-01	0.010367	-0.002486	0.006926	0.004399	0.009482	0.025854	
2006-03-31	0.014995	0.001925	-0.002580	0.008029	0.026566	0.039604	
2006-04-30	-0.040435	0.000384	-0.023276	-0.039826	-0.029785	-0.057143	

	B17517	B05851	B12869	B20167	B18365	B21754	\
2005-12-31	0.055187	0.002578	0.056000	-0.128571	0.029695	6.471691e-04	
2006-01-30	0.005231	0.004005	0.001894	0.082367	0.008614	2.882650e-03	
2006-03-01	0.030647	-0.006759	0.017013	-0.023642	0.003713	1.107340e-03	
2006-03-31	0.049696	-0.000390	0.015799	-0.046916	-0.010359	2.414173e-03	
2006-04-30	-0.044445	0.003683	-0.046661	0.056372	-0.032150	3.509704e-07	

	B21401	B21937	B21046	B18646	B11920	B20159	\
2005-12-31	0.038283	0.020202	0.054011	0.067529	0.013793	0.028766	
2006-01-30	-0.001034	0.015993	-0.002030	-0.008274	0.011346	0.000508	
2006-03-01	0.004484	0.004843	0.036096	0.031491	-0.009879	0.071138	
2006-03-31	0.028159	0.000747	0.009323	0.005258	0.008375	0.055503	
2006-04-30	-0.029726	-0.034328	-0.057365	-0.057562	-0.017026	0.004045	

	B07325	B21967
2005-12-31	0.060254	0.131593
2006-01-30	0.012961	0.011300
2006-03-01	0.036910	-0.000466
2006-03-31	0.046511	0.074523
2006-04-30	-0.043991	-0.114434


```
[6]: for method in ['bootstrap', 'historical', 'normal', 't-dist']:
      simOut = calcVaR(dfrtn, method=method)
      VaR, CVaR = simOut['VaR'], simOut['CVaR']
      print(f'method = {method}, VaR = {VaR}, CVaR = {CVaR}')
```

100%| | 100000/100000 [00:02<00:00, 35924.79it/s]

```
method = bootstrap, VaR = 0.09459406695868117, CVaR = 0.13663232164504419
method = historical, VaR = 0.1549411556439678, CVaR = 0.16770678995134303
method = normal, VaR = 0.07019531506080463, CVaR = 0.08120307869127266
method = t-dist, VaR = 0.09625891023369396, CVaR = 0.18400253654661636
```

```
[7]: portRtn = simOut['portRtn']
      print(f'portfolio return mean = {np.mean(portRtn)}, sd = {np.std(portRtn)}')
      plt.hist(portRtn, color='k', bins=40, density=True)
      plt.xlabel('portfolio return')
      plt.ylabel('density')
      plt.show()
```

portfolio return mean = 0.005373969994835999, sd = 0.03248408627914174



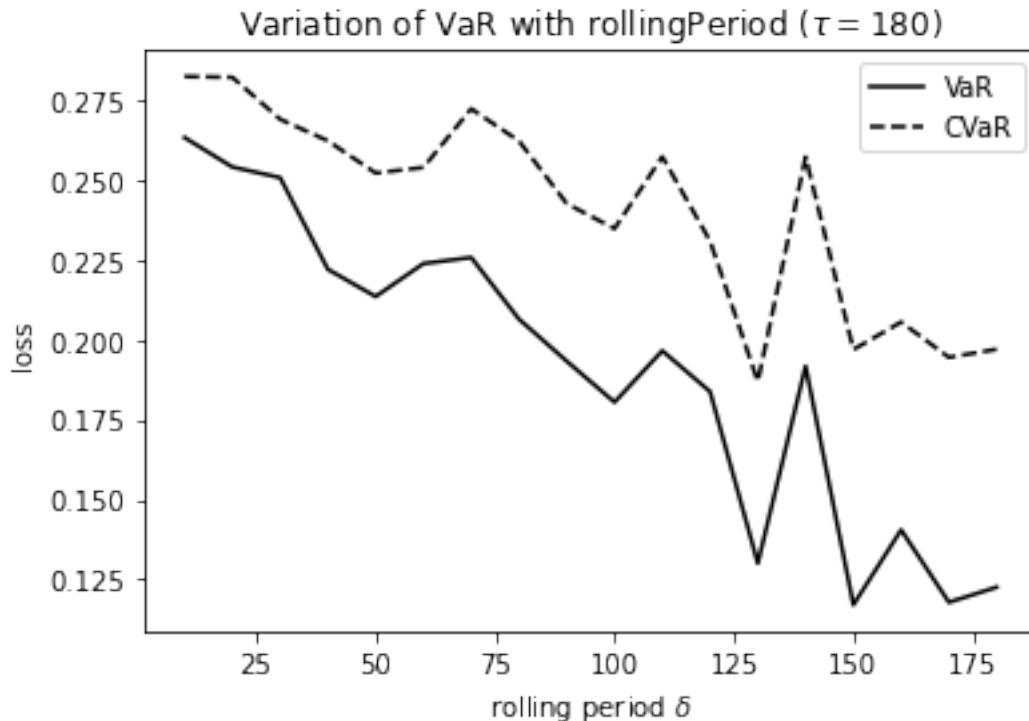
5.1.1 Variation of VaR with rollingPeriod

Here we study the variation of VaR with rollingPeriod δ . When δ is small, returns have large overlapping window, thus more correlated, which leads to more severe VaR. Intuition: let's say there is a market crash, when we have a small δ , impact of the crash persists for a longer period of

time (i.e. reflected in our computed returns), thus low returns stand higher chance getting sampled, leading to large VaR. Below, we use $\delta = 10, 20, \dots, 180$ keeping `timeHorizon` fixed at 180 (half-year). Indeed VaR and CVaR drop and stabilize with larger δ , when overlapping is less.

```
[8]: rollingPeriods = np.arange(10,190,10)
varList = list()
cvarList = list()
for rollingPeriod in rollingPeriods:
    dfrtn = calcHistoricalReturns(df, rollingPeriod, 180)
    simOut = calcVaR(dfrtn, method='bootstrap', Nsim=1e4)
    varList.append(simOut['VaR'])
    cvarList.append(simOut['CVaR'])
plt.plot(rollingPeriods,varList,'k',label='VaR')
plt.plot(rollingPeriods,cvarList,'k--',label='CVaR')
plt.title(r'Variation of VaR with rollingPeriod ($\tau=180$)')
plt.xlabel('rolling period $\delta$')
plt.ylabel('loss')
plt.legend()
plt.show()
```

```
100%|      | 10000/10000 [00:00<00:00, 19653.39it/s]
100%|      | 10000/10000 [00:00<00:00, 28895.55it/s]
100%|      | 10000/10000 [00:00<00:00, 36028.12it/s]
100%|      | 10000/10000 [00:00<00:00, 37559.68it/s]
100%|      | 10000/10000 [00:00<00:00, 42261.18it/s]
100%|      | 10000/10000 [00:00<00:00, 43059.49it/s]
100%|      | 10000/10000 [00:00<00:00, 43452.94it/s]
100%|      | 10000/10000 [00:00<00:00, 43997.04it/s]
100%|      | 10000/10000 [00:00<00:00, 46278.45it/s]
100%|      | 10000/10000 [00:00<00:00, 46188.91it/s]
100%|      | 10000/10000 [00:00<00:00, 46362.65it/s]
100%|      | 10000/10000 [00:00<00:00, 46555.58it/s]
100%|      | 10000/10000 [00:00<00:00, 46763.78it/s]
100%|      | 10000/10000 [00:00<00:00, 46988.95it/s]
100%|      | 10000/10000 [00:00<00:00, 46617.41it/s]
100%|      | 10000/10000 [00:00<00:00, 46864.36it/s]
100%|      | 10000/10000 [00:00<00:00, 46950.55it/s]
100%|      | 10000/10000 [00:00<00:00, 48065.15it/s]
```



5.1.2 Variation of VaR with start/endDate


Here we study the variation of VaR with **start/endDate** $t_{s/e}$, fixing **rollingPeriod** at 7 and **timeHorizon** at 7, i.e. non-overlapping weekly returns. The routine offers flexibility for user to choose different start and end dates, to inspect periods of interest. Historical returns between **startDate** and **endDate**, spanning one year, are fed into VaR routine and we examine the VaR time series. As expected, during financial crisis in 2008, assets do not perform well, which is reflected in the spike in VaR and CVaR; while during the market boom starting 2010 onwards, VaR has relatively low values, of the order 3%.

```
[9]: N = 12
T0 = pd.date_range(start='01-01-2006', periods=N, freq='1Y')
T1 = T0 + pd.DateOffset(years=1)
varList = list()
cvarList = list()
for i in range(N):
    dfrtn = calcHistoricalReturns(df, 7, 7, T0[i], T1[i])
    simOut = calcVaR(dfrtn, method='bootstrap', Nsim=1e4)
    varList.append(simOut['VaR'])
    cvarList.append(simOut['CVaR'])
plt.plot(T0, varList, 'k', label='VaR')
plt.plot(T0, cvarList, 'k--', label='CVaR')
plt.title(r'Variation of VaR with start/endDate ($\tau=7$)')
```



```
plt.xlabel('start date $t_s$')
plt.ylabel('loss')
plt.legend()
plt.show()
```

```
100%|      | 10000/10000 [00:00<00:00, 46220.52it/s]
100%|      | 10000/10000 [00:00<00:00, 46352.92it/s]
100%|      | 10000/10000 [00:00<00:00, 46346.05it/s]
100%|      | 10000/10000 [00:00<00:00, 46501.44it/s]
100%|      | 10000/10000 [00:00<00:00, 46273.81it/s]
100%|      | 10000/10000 [00:00<00:00, 46331.26it/s]
100%|      | 10000/10000 [00:00<00:00, 46070.45it/s]
100%|      | 10000/10000 [00:00<00:00, 46562.14it/s]
100%|      | 10000/10000 [00:00<00:00, 46405.64it/s]
100%|      | 10000/10000 [00:00<00:00, 46413.85it/s]
100%|      | 10000/10000 [00:00<00:00, 46323.53it/s]
100%|      | 10000/10000 [00:00<00:00, 46566.90it/s]
```



5.2 Bootstrap mutual fund paths

Here we follow a different bootstrap approach, known as **historical bootstrap**, more commonly used in risk management:

1. bootstrap from historical **daily** mutual fund returns to simulate portfolio NAV over the next

`timeHorizon`

2. repeat bootstrap procedure for `Nsim` scenarios, from which VaR and CVaR are calculated

The difference with previous approach is, we simulate entire paths of mutual funds, day by day, for a future `timeHorizon`, and then we have a collection of future scenarios from which VaR and CVaR can be estimated. `histBootVaR(df, ...)` handles this routine. As very old data may be less relevant, parameter `lookback` is included, for example, `lookback=90` uses only the most recent 90 days of data for bootstrap.

Below we study calculations bootstrapping from the entire historical timeframe and recent 90 days. As recent 90 days do not involve extreme market crash, VaR and CVaR for `lookback=90` is smaller, by roughly a half. Compared with the previous approach (bootstrap portfolio returns), this approach gives smaller VaR estimates, as we are simulating day by day, and daily gains and losses even out, leading to small VaR.

```
[10]: simOut = histBootVaR(df, 30)
      VaR, CVaR = simOut['VaR'], simOut['CVaR']
      print(f'method = historical bootstrap, VaR = {VaR}, CVaR = {CVaR}')
```

```
100%|          | 100000/100000 [00:39<00:00, 2550.01it/s]
```

```
method = historical bootstrap, VaR = 0.07972623738332185, CVaR =
0.09354835204704694
```

```
[11]: simOut = histBootVaR(df, 30, lookback=90)
      VaR, CVaR = simOut['VaR'], simOut['CVaR']
      print(f'method = historical bootstrap, VaR = {VaR}, CVaR = {CVaR}')
```

```
100%|          | 100000/100000 [00:39<00:00, 2562.45it/s]
```

```
method = historical bootstrap, VaR = 0.03634564642711302, CVaR =
0.04262776968178128
```